

MITESH ADAKE

Plano, TX | +1 (901) 651-7876 | miteshadake29@gmail.com | [linkedin.com/in/mitesh-adake](https://www.linkedin.com/in/mitesh-adake) | mitadake.github.io/

EDUCATION

University of Southern California

Master's, Computer Science

• GPA: 3.75/ 4.0

• Coursework: Applied Natural Language Processing, Advanced Computer Vision, Adversarial and Trustworthy Foundation Models

Aug 2024 - May 2026

Los Angeles, CA, USA

Pune Institute of Computer Technology

Bachelor's, Computer Engineering

• GPA: 3.84/ 4.0

• Coursework: Data Structures & Algorithms, Machine Learning, Natural Language Processing, Deep Learning

Jul 2019 - May 2023

Pune, MH, India

SKILLS

- **Programming Languages:** Python, Java
- **Databases / Query Languages:** MongoDB, MySQL, Oracle (SQL)
- **Cloud Platforms:** Amazon Web Services, Google Cloud Platform
- **AI / Machine Learning:** Neural Networks, Natural Language Processing, Large Language Models, Transformers, PyTorch, Tensorflow, Keras, Model Context Protocol, LangGraph
- **Tools & Frameworks:** Ab Initio, Dataiku, Docker, Flask, FastAPI, Django, Power BI

PROFESSIONAL EXPERIENCE

Federal Express Corporation

Sr. Revenue Science Analyst

Plano, TX, USA · Jun 2026 - Present

- Driving AI initiatives integrating AI agents, RAG, LLMs, and ML models with LangGraph, Model Context Protocol (MCP), and enterprise tools to automate pricing workflows and enhance operational efficiency for Pricing Analysts and Sales teams.

Revenue Management - Intern

Los Angeles, CA, USA · Sep 2025 - May 2026

- Built 3 automated systems leveraging data engineering to automate pricing analyst workflows, cutting processing time from ~5 days to ~1 hour
- Partnered with cross-functional teams on PowerBI to create dashboards and AI proof-of-concepts for data-driven pricing systems

Revenue Management - Summer Intern

Memphis, TN, USA · Jun 2025 - Aug 2025

- Designed and developed an automated system using Python, SQL, Streamlit, and Docker to process user inputs and perform complex, business-critical database operations
- Reduced manual processing time from ~10 days to ~10 minutes to optimize workflow logic, significantly increasing operational efficiency

Software Engineer

Hyderabad, TL, India · Aug 2023 - Jul 2024

- Migrated Dataiku data pipeline to Ab Initio (ETL) data pipeline for Harmonized Code search engine, resulting in annual cost savings of \$ 1M for FedEx and accelerating data processing time
- Executed an optimization strategy in Python for Natural Language Processing (NLP) code, significantly reducing data discrepancies from 25% to negligible levels and eliminating data shuffling from 75% of text data

Persistent Systems

Machine Learning Intern

Pune, MH, India · Jun 2022 - Aug 2022

- Implemented NLP-based question deduplication on company's community portal, improving efficiency and reducing duplicate postings to 10%
- Leveraged Python, Flask API, FAST API, and Docker to create a scalable and robust solution

NICE

Data Engineer Intern

Pune, MH, India · Nov 2021 - Apr 2022

- Built scalable, serverless real-time file ingestion using AWS (S3, DynamoDB, Lambda) and MongoDB for efficient file processing
- Streamlined data management tasks by implementing a user-friendly Python Flask interface

PROJECTS & OUTSIDE EXPERIENCE

Fast-Hack: GGUF k-quant Content-Injection Attack [Project Link](#)

Jan 2026 - May 2026

- Researched and built a new and faster quantization attack technique for LLMs: stay benign under full-precision (FP) inference but become malicious after GGUF k-quantization to produce content injected output
- Evaluated on Llama-3.2-1B-Instruct and Qwen2.5-1.5B-Instruct across Q2_K, Q3_K_M, Q4_K_M, Q5_K_M, and Q6_K with ASR up to ~82–83% on Q4_K_M with FP ASR ~0–3%; other quants showed ~38–68% quantized ASR with ~2–18% FP

Hate Speech Unlearning for LLMs

Jan 2025 - Mar 2025

- Implemented Task Vector Arithmetic, Contrastive Learning, and Fine-tuning to remove hate speech from LLaMA 3.1 models, resulting in a drop of harmful rate from 62% to 34% on PKU-SafeRLHF dataset

PriceNet: Stock Price Prediction using Large Language Models

Sep 2024 - Nov 2024

- Designed and implemented PriceNet, using Llama and Gemini for explainable financial time series forecasting with accuracy of 50% and ROGUE-2 of 0.546 on reasoning of predictions